



Scaling Attributed Network Embedding to Massive Graphs

Renchi Yang, Jieming Shi, Xiaokui Xiao,
Yin Yang, Juncheng Liu, Sourav S. Bhowmick





Outline

- Problem Definition & Applications
- Existing Work & Limitations
- Basic Idea & Challenges
- Random Walks & Affinity Measure
- Objective Function & Solution
- Parallel Implementations
- Experiments

Attributed Network

- **Input:** a graph G , where each node u has some attributes from a set R , e.g., a social network user has an attribute “gender”




Mark Zuckerberg

| | |
|--------------|----------|
| Gender: M | M: 1 |
| Country: USA | F: 0 |
| Age: 37 | Other: 0 |
| #fans: 116M | |

- Each node-attribute pair (u, r) has a weight $w(u, r)$ that indicates the strength of association


Attributed Network Embedding (ANE)

- **Objective:** Map each node to an embedding vector, which can then be used as input to downstream machine learning tasks



Applications

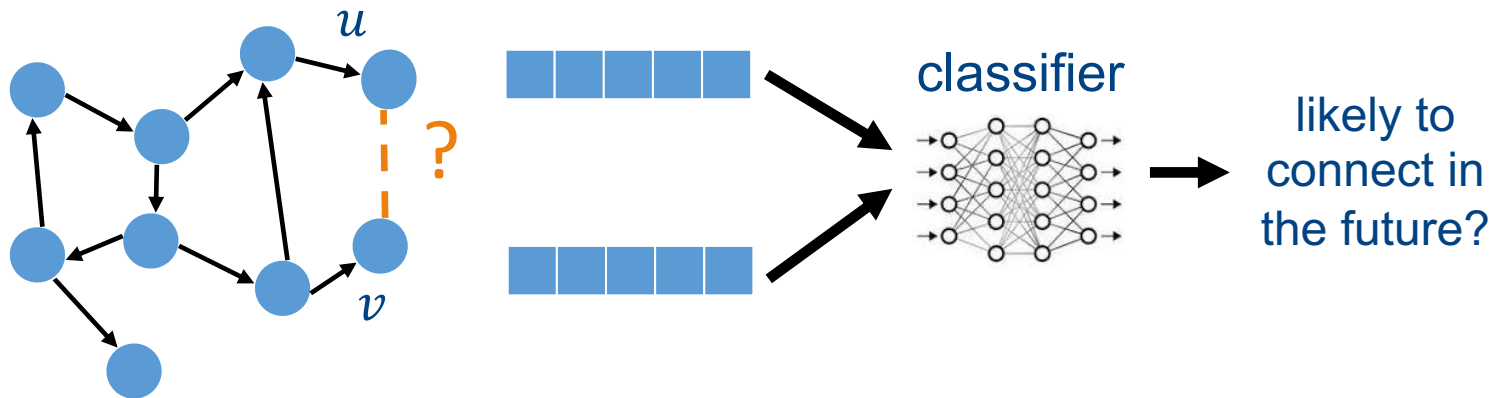
- Node classification
 - user tagging in social networks
 - fraud detection in financial networks
 - cancer biomarkers identification in biological networks



Applications

- Link Prediction

- image/video recommendation in Pinterest [Ying KDD'18]
- product recommendation in Alibaba [Cen KDD'19]

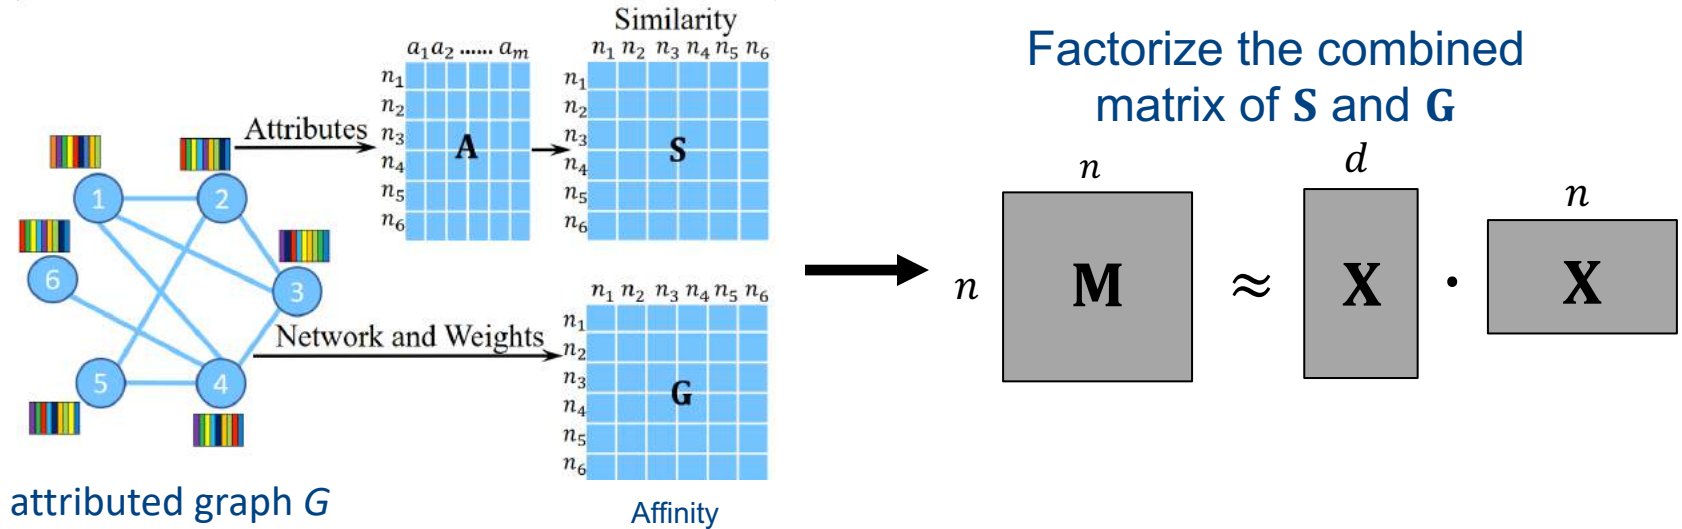


Ying et al. Graph Convolutional Neural Networks for Web-Scale Recommender Systems. KDD'18.

Cen et al. Representation Learning for Attributed Multiplex Heterogeneous Network. KDD'19.

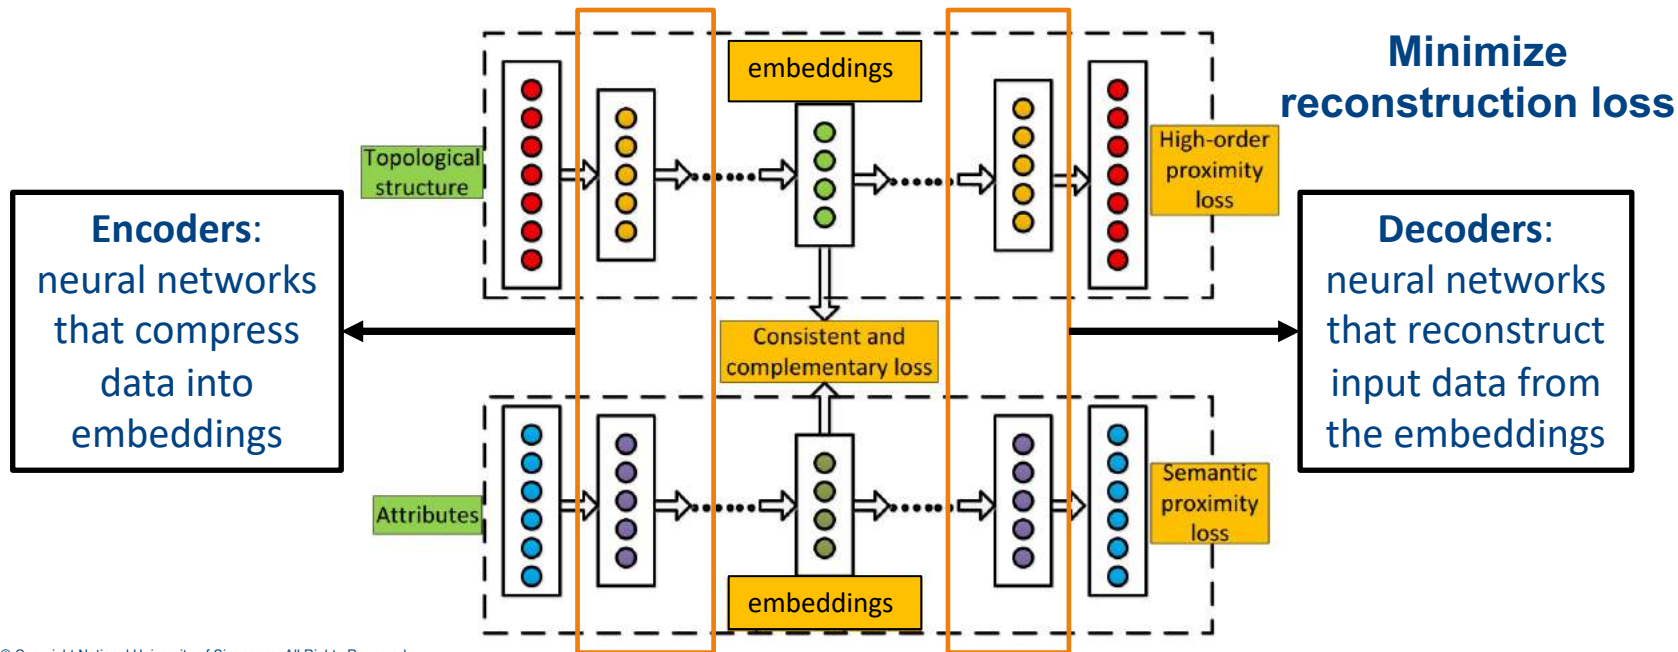
Existing work

- Matrix Factorization-based methods




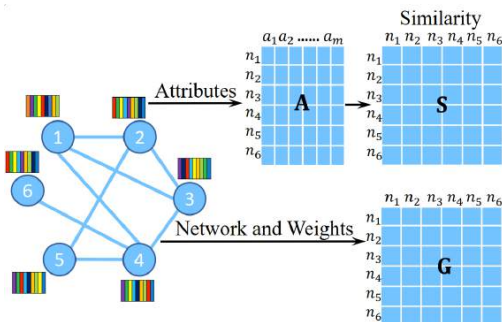
Existing work

- Neural networks-based methods




Limitations

- Existing ANE solutions
 - generate effective embeddings based on expensive learning techniques (e.g., auto-encoders). This makes them **difficult to handle large graphs**.
 - are more efficient but produce **low-quality embeddings**.






Basic idea & challenges

- Objective: Construct high-quality embeddings without significant computation cost
- Basic idea: formulate ANE as a new matrix factorization problem
- Challenges
 - Two types of affinity to capture: node-node & node-attribute
 - How to model these affinities?
 - How to compute & store these affinities?
($O(n^2)$ space, $n=\text{\#nodes}$)
 - Two affinity matrices
 - Which one are we factorizing?
 - How to design the objective function?




Solution Overview: PANE



- Construct node-attribute affinity matrix \mathbf{F} & attribute-node affinity matrix \mathbf{B}
 - forward & backward random walk models
 - indirectly model node-node affinity via \mathbf{F} & \mathbf{B}
- Joint factorization of affinity matrices \mathbf{F} & \mathbf{B}
 - initialize embeddings via singular value decomposition
- Parallelize PANE for higher efficiency
 - parallel singular value decomposition



Two Types of Random Walks




Forward: node-to-attribute



Backward: attribute-to-node

Forward Random Walks


- Forward random walk from node u :
 - Start from u
 - At each step, stop with α probability
 - After stopping at a node v , pick an attribute r with probability $\propto w(v, r)$
- Intuition: it samples an attribute r from the vicinity of u
- $p_f(u, r)$ is the probability that a forward random walk from u samples r in the end



Node-Attribute Affinity

- Our node-attribute affinity measure:

$$\mathbf{F}[u, r] = \log \left(\frac{n \cdot p_f(u, r)}{\sum_{v \in V} p_f(v, r)} + 1 \right)$$




- Understand $\mathbf{F}[u, r]$ using PMI (pointwise mutual information) in information theory:
 - PMI $\log\left(\frac{\Pr(r|v)}{\Pr(r)}\right)$ measures the co-occurrence of elements and in a collection of elements S
 - all random walks as S , $\sum_{v \in V} p_f(v, r)$ as $\Pr(r)$ and $p_f(u, r)$ as $\Pr(r|v)$
 - $\mathbf{F}[u, r]$ measures how frequently u, r co-occur on all random walks to r
 - Word2vec is implicitly factorizing a PMI matrix to obtain word embeddings [Levy NeurIPS'14]

Computing node-attribute affinity

- Node-attribute affinity

- $$\mathbf{F}[u, r] = \log \left(\frac{n \cdot p_f(u, r)}{\sum_{v \in V} p_f(v, r)} + 1 \right)$$
- $$p_f(u, r) = \alpha \sum_{i=0}^t (1 - \alpha)^i \mathbf{P}^i \cdot \mathbf{R}_r[u, r]$$
- $O(mdt)$ time using power method




| | v_1 | v_2 | v_3 | v_4 |
|-------|-------|---------------|---------------|---------------|
| v_1 | 0 | $\frac{1}{3}$ | $\frac{1}{3}$ | $\frac{1}{3}$ |
| v_2 | 0 | 0 | 1 | 0 |
| v_3 | 0 | 0 | 0 | 1 |
| v_4 | 0 | 0 | 1 | 0 |

Transition matrix \mathbf{P}

| | r_1 | r_2 | r_3 |
|-------|---------------|---------------|---------------|
| v_1 | $\frac{1}{2}$ | $\frac{1}{2}$ | 0 |
| v_2 | $\frac{1}{3}$ | 0 | $\frac{2}{3}$ |
| v_3 | $\frac{1}{3}$ | $\frac{1}{3}$ | $\frac{1}{3}$ |
| v_4 | 1 | 0 | 0 |

Row-normalized attribute matrix \mathbf{R}_r




Backward RW & Attribute-Node Affinity

- Backward random walk from attribute r
 - Randomly pick a node s with probability $\propto w(s, r)$
 - Start a random walk from s
 - At each step, stop with α probability
 - Let v be the stopping point of the walk
- Our attribute-to-node affinity measure:

$$\mathbf{B}[r, v] = \log \left(\frac{d \cdot p_b(r, v)}{\sum_{r' \in R} p_b(r', v)} + 1 \right)$$

- where R is the set of all attributes,
- $p_b(r, v)$ is the probability that a backward random walk from r samples v in the end



Why node-attribute & attribute-node affinities

- We can indirectly model node-node affinities with much less space using node-attribute + attribute-node affinities



- Intuition:
 - Consider a forward random walk from u to r & a backward random walk from r to v
 - They form an extended walk from u to v , in which we "teleport" through a virtual connection by r
 - Such random walks could be combined to model node-node affinity

Capturing node-node affinity

- Our node-node affinity can be indirectly constructed via:


$$p(u, v) = \sum_{r \in R} \mathbf{F}[u, r] \cdot \mathbf{B}[r, v]$$



- Thus, we do not need an $n \times n$ node-node affinity matrix explicitly
 - space overhead: $O(n^2) \rightarrow O(nd), d \ll n$
 - $d = \text{\#attributes}, n = \text{\#nodes}$

Objective function

- We construct
 - two embedding matrices $\mathbf{X}_f, \mathbf{X}_b \in \mathbb{R}^{n \times k}$ for the nodes, and
 - one embedding matrix $\mathbf{Y} \in \mathbb{R}^{d \times k}$ for attributes
- Optimization objective:
 - $\min_{\mathbf{X}_f, \mathbf{Y}, \mathbf{X}_b} \|\mathbf{F} - \mathbf{X}_f \mathbf{Y}^T\|_F^2 + \|\mathbf{B} - \mathbf{X}_b \mathbf{Y}^T\|_F^2$
 - $\mathbf{X}_f \cdot \mathbf{Y}^T \approx \mathbf{F}$, to capture node-attribute affinity
 - $\mathbf{X}_b \mathbf{Y}^T \approx \mathbf{B}$, to capture attribute-node affinity




Solving the optimization objective

- Optimization objective:

$$\min_{\mathbf{X}_f, \mathbf{Y}, \mathbf{X}_b} \|\mathbf{F} - \mathbf{X}_f \mathbf{Y}^T\|_F^2 + \|\mathbf{B} - \mathbf{Y} \cdot \mathbf{X}_b^T\|_F^2$$


- We can obtain \mathbf{X}_f , \mathbf{X}_b , and \mathbf{Y} using gradient descent



- a large number of iterations are required till convergence
- jointly updating \mathbf{X}_f , \mathbf{X}_b and \mathbf{Y} involves intensive computation

- Our idea for efficiency:

- find a good initialization for \mathbf{X}_f , \mathbf{X}_b and \mathbf{Y} based on singular value decomposition (SVD)




Initialize embeddings via SVD


- Initial solution for the first part via randomized SVD

$$\mathbf{F} \approx \mathbf{U} \cdot \mathbf{\Sigma} \cdot \mathbf{V}^T$$

$\mathbf{X}_f \leftarrow \mathbf{U} \quad \mathbf{V}^T \rightarrow \mathbf{Y}$




- Initial solution for the second part
 - approximate singular vectors $\mathbf{Y} = \mathbf{V}$ is semi-unitary, i.e., $\mathbf{Y}^T \mathbf{Y} = \mathbf{I}$
 - Intuitively, if we want $\mathbf{X}_b \mathbf{Y}^T = \mathbf{B}$,
 $\mathbf{X}_b = \mathbf{X}_b \mathbf{Y}^T \mathbf{Y} = \mathbf{B} \cdot \mathbf{Y}$




Space & time complexities

- Time complexity:
 - $O(mdt + ndkt)$
 - $t = \text{\#iterations of gradient descent}$
($t = 5$ in our experiments)
 - $k = \text{the embedding size}$
 - $m = \text{\#edges}$
 - $n = \text{\#nodes}$
 - $d = \text{\#attributes}$
- Space complexity
 - $O(m + nd + nk)$



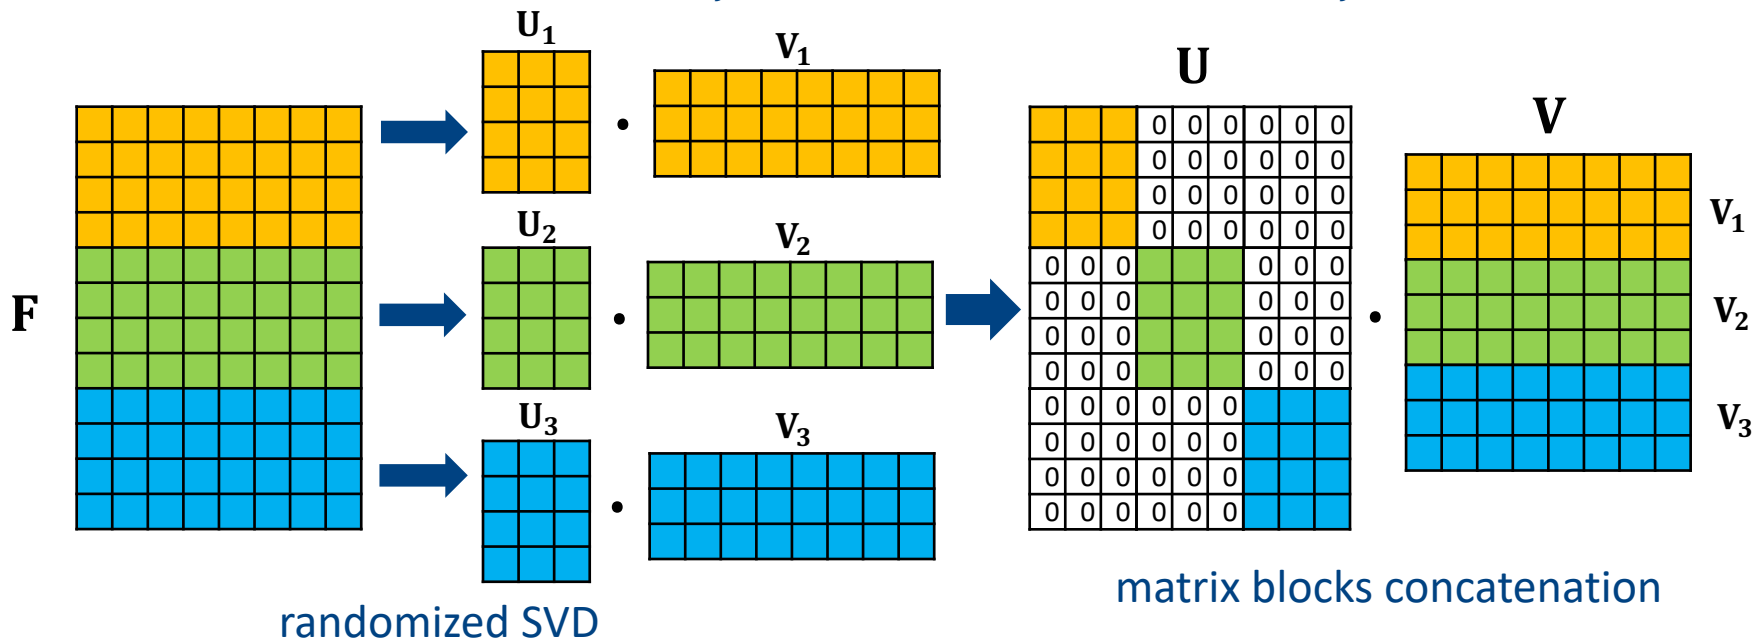
Parallel implementation of PANE

- Parallel computation of \mathbf{F} and \mathbf{B}
 - Matrix multiplications can be easily parallelized
- Parallel gradient descent. In each iteration,
 - $\mathbf{X}_f[u] \forall u \in V$ or $\mathbf{X}_b[v] \forall v \in V$ can be updated independently
 - $\mathbf{Y}[r]$ for each attribute r can be updated independently
- Parallel SVD over \mathbf{F}
 - The SVD over $\mathbf{F} \neq$ direct SVD over each matrix blocks of \mathbf{F}
 - How to perform SVD over \mathbf{F} in parallel?



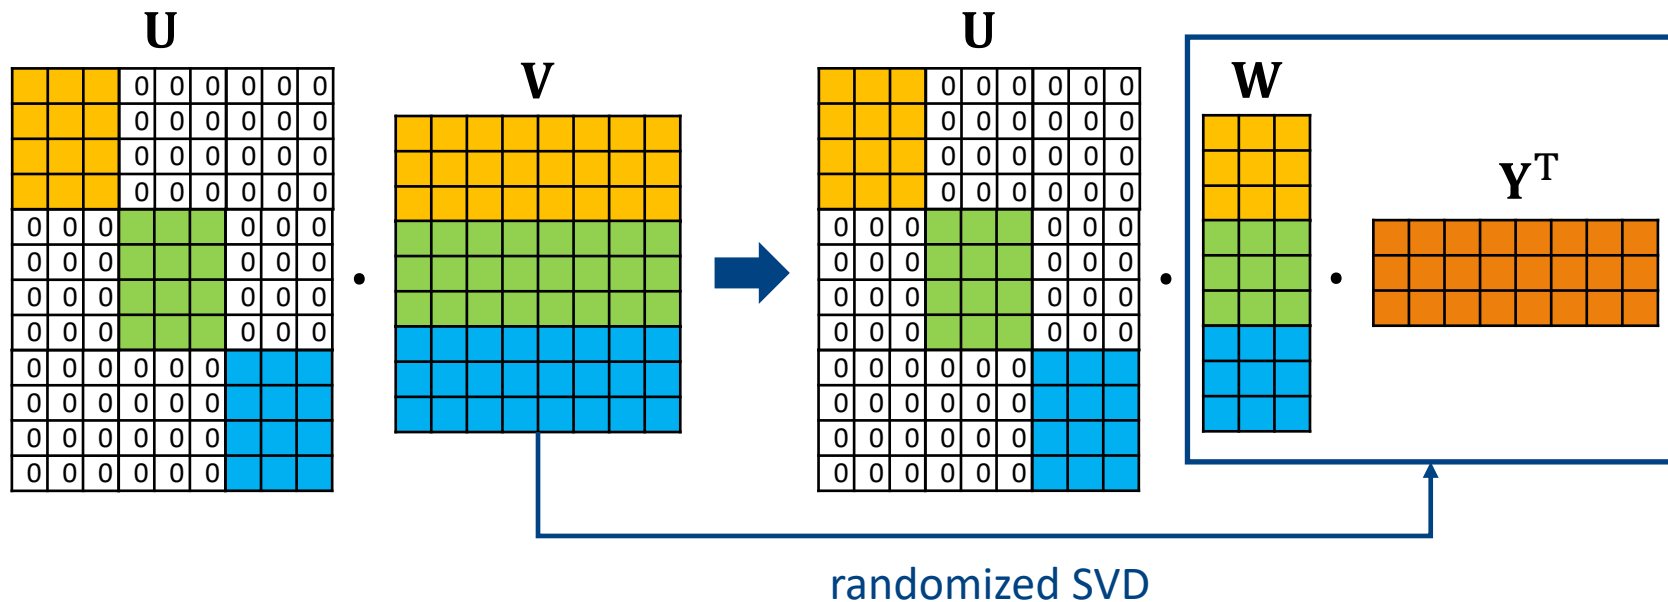
Parallel SVD

- Given $\mathbf{F} \in \mathbb{R}^{12 \times 8}$, we want $\mathbf{X}_f \in \mathbb{R}^{12 \times 3}$ and $\mathbf{Y} \in \mathbb{R}^{8 \times 3}$ s.t. $\mathbf{X}_f \cdot \mathbf{Y}^T \approx \mathbf{F}$



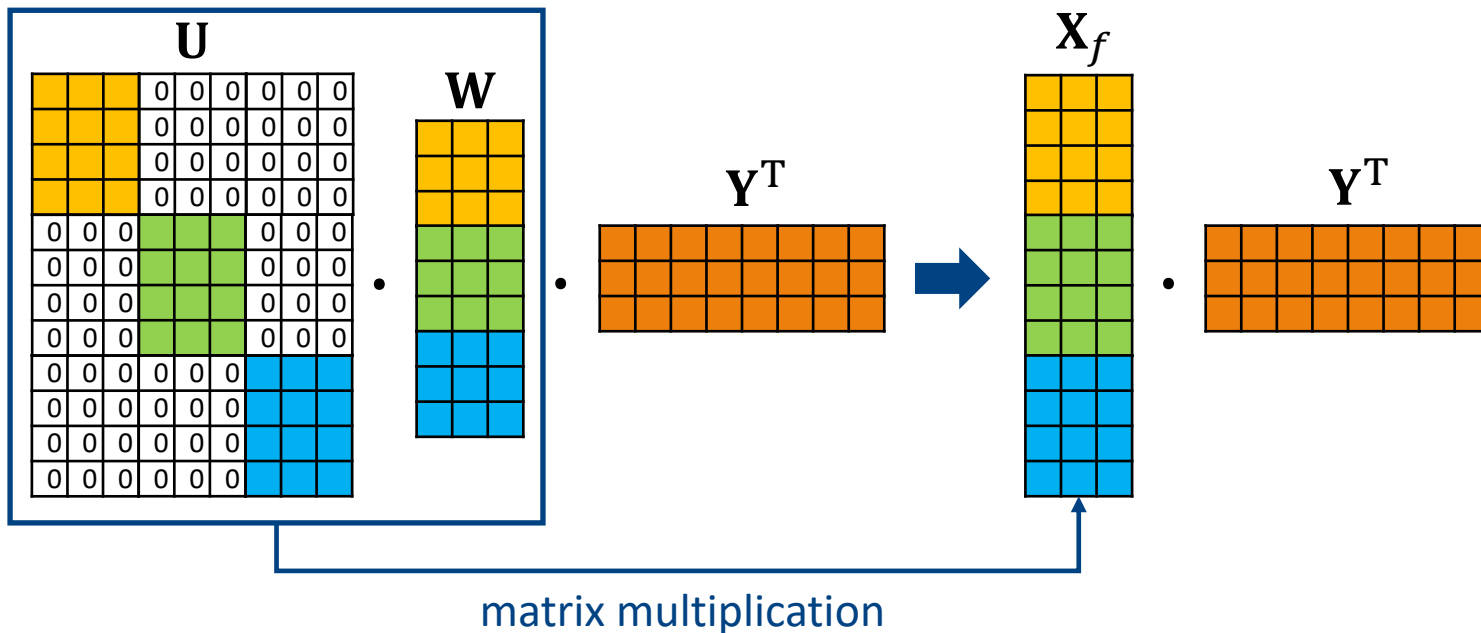
Parallel SVD

- Given $\mathbf{F} \in \mathbb{R}^{12 \times 8}$, we want $\mathbf{X}_f \in \mathbb{R}^{12 \times 3}$ and $\mathbf{Y} \in \mathbb{R}^{8 \times 3}$ s.t. $\mathbf{X}_f \cdot \mathbf{Y}^T \approx \mathbf{F}$




Parallel SVD

- Given $\mathbf{F} \in \mathbb{R}^{12 \times 8}$, we want $\mathbf{X}_f \in \mathbb{R}^{12 \times 3}$ and $\mathbf{Y} \in \mathbb{R}^{8 \times 3}$ s.t. $\mathbf{X}_f \cdot \mathbf{Y}^T \approx \mathbf{F}$



Parallel SVD

- Given $\mathbf{F} \in \mathbb{R}^{12 \times 8}$, we want $\mathbf{X}_f \in \mathbb{R}^{12 \times 3}$ and $\mathbf{Y} \in \mathbb{R}^{8 \times 3}$ s.t. $\mathbf{X}_f \cdot \mathbf{Y}^T \approx \mathbf{F}$



Experiments: Datasets

| Name | # of nodes | # of edges | # of distinct attributes | # of attributes per node | # of distinct labels |
|----------|------------|------------|--------------------------|--------------------------|----------------------|
| Cora | 2.7k | 5.4k | 1.4k | 18.2 | 7 |
| Citeseer | 3.3k | 4.7k | 3.7k | 31.9 | 6 |
| Facebook | 4k | 88.2k | 1.3k | 8.3 | 193 |
| Pubmed | 19.7k | 44.3k | 0.5k | 50.2 | 3 |
| Flickr | 7.6k | 479.5k | 12.1k | 24.0 | 9 |
| Google+ | 107.6k | 13.7M | 15.9k | 2793.7 | 468 |
| TWeibo | 2.3M | 50.7M | 1.7k | 7.3 | 8 |
| MAG | 59.3M | 978.2M | 2.0k | 7.3 | 100 |

Experiments: Competitors

- PANE: Our solution for attributed graphs
- Default embedding dimensionality: $k = 128$
- CPU: Intel Xeon 2.2GHz

■ 6 neural-based methods

- STNE [KDD 2018]
- ARGAS [IJCAI 2018]
- LQANR [IJCAI 2019]
- CAN [WSDM 2019]
- DGI [ICLR 2019]
- GATNE [KDD 2019]

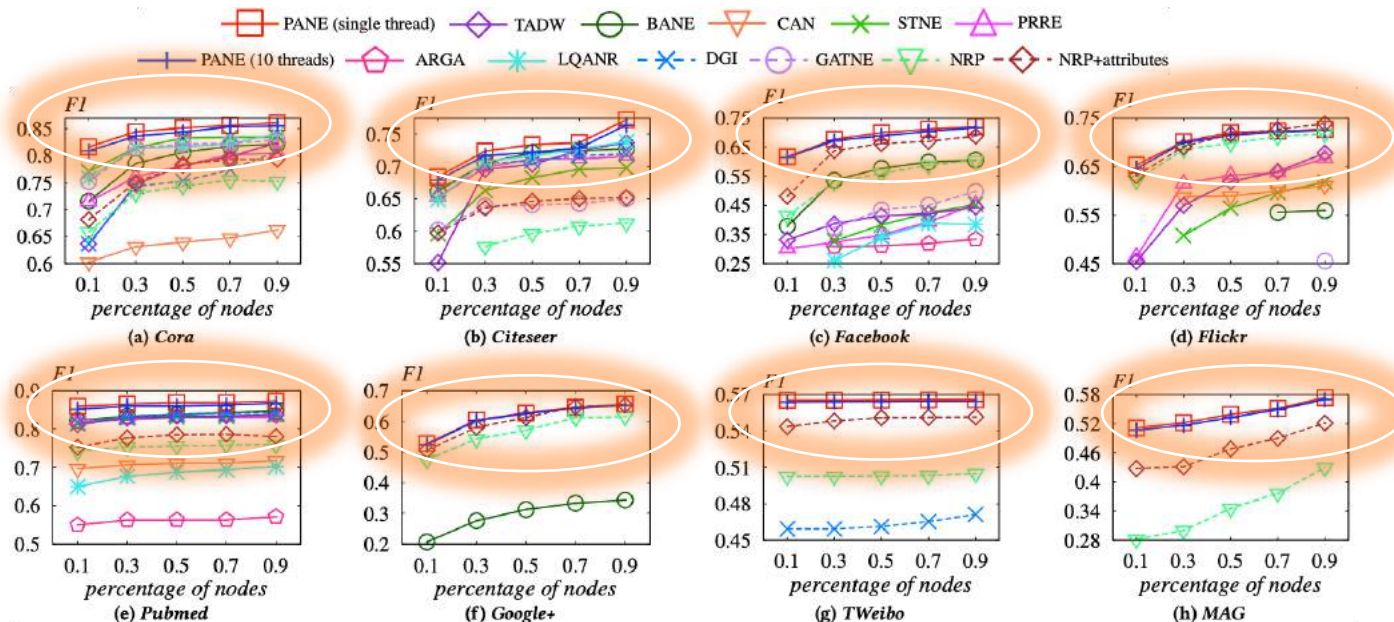
■ 3 factorization-based methods

- TADW [IJCAI 2015]
- BANE [ICDM 2018]
- NRP [VLDB 2020]

■ 1 other method

- PRRE [CIKM 2018]

Results: Node Classification

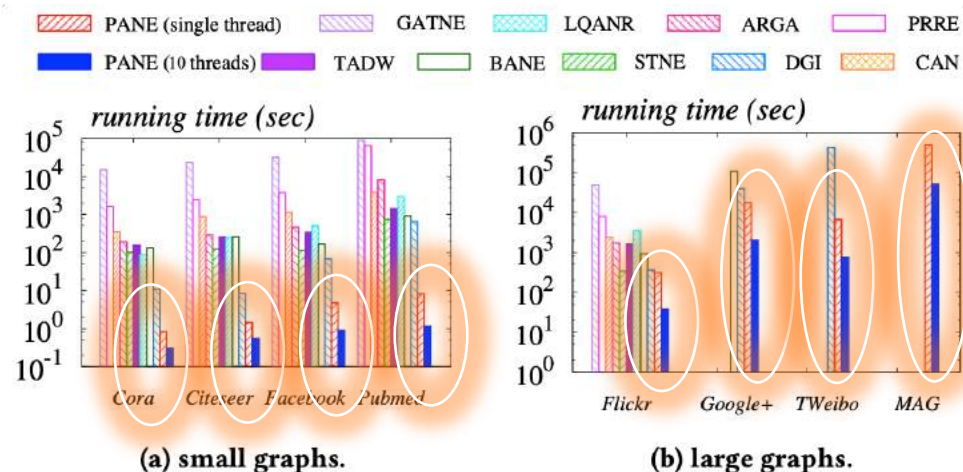


- Percentage of nodes used for training: 10% ~ 90%
- PANE vs. SOTA: improvements of 3.4%-17.2% in terms of F1 measure

Results: Link Prediction

| Method | Area Under Curve (AUC) | | | | | | | |
|----------------------|------------------------|-----------------|---------------|-----------------|---------------|----------------|---------------|------------|
| | <i>Cora</i> | <i>Citeseer</i> | <i>Pubmed</i> | <i>Facebook</i> | <i>Flickr</i> | <i>Google+</i> | <i>TWeibo</i> | <i>MAG</i> |
| NRP | 0.796 | 0.86 | 0.87 | 0.969 | 0.909 | 0.989 | 0.967 | 0.915 |
| GATNE | 0.791 | 0.687 | 0.745 | 0.961 | 0.805 | - | - | - |
| TADW | 0.829 | 0.895 | 0.904 | 0.752 | 0.573 | - | - | - |
| ARGA | 0.64 | 0.637 | 0.623 | 0.71 | 0.676 | - | - | - |
| BANE | 0.875 | 0.899 | 0.919 | 0.796 | 0.64 | 0.56 | - | - |
| PRRE | 0.879 | 0.895 | 0.887 | 0.899 | 0.789 | - | - | - |
| STNE | 0.808 | 0.71 | 0.789 | 0.962 | 0.638 | - | - | - |
| CAN | 0.663 | 0.734 | 0.734 | 0.714 | 0.5 | - | - | - |
| DGI | 0.51 | 0.5 | 0.73 | 0.711 | 0.769 | 0.792 | 0.721 | - |
| LQANR | 0.886 | 0.916 | 0.904 | 0.951 | 0.824 | - | - | - |
| PANE (single thread) | 0.933 | 0.932 | 0.985 | 0.982 | 0.929 | 0.987 | 0.976 | 0.96 |
| PANE (10 threads) | 0.929 | 0.929 | 0.985 | 0.98 | 0.927 | 0.984 | 0.975 | 0.958 |

Results: Efficiency




- Compared to the state of the art, PANE is **orders of magnitude faster**
- On the MAG dataset with 0.98 billion edges, PANE can terminate within 12 hours using 10 CPU cores (Intel Xeon 2.2GHz)


Effectiveness of SVD-based initializations

PANE-R: the algorithm that uses random initializations


PANE: the algorithm that uses SVD-based initializations



(a) Facebook.





(b) Pubmed.



(c) Flickr.

Link prediction performance vs. running time
when varying #iteration for the gradient descent from 1 to 20

THANK YOU




$$\begin{matrix} \mathbf{X}_f & \cdot & \mathbf{Y}^T & \approx & \mathbf{F}[u, r] \\ \text{for } u & & \text{for } r & & \end{matrix}$$

$$\begin{matrix} \mathbf{X}_b & \cdot & \mathbf{Y}^T & \approx & \mathbf{B}[u, r] \\ \text{for } u & & \text{for } r & & \end{matrix}$$




1

Random walks



2

Joint matrix factorization



3

Parallelization